

A Programmatic Interface for Particle Plasma Simulation in Python

PRELIMINARY RESULTS WITH PYCUDA



 python
Min Ragan-Kelley (UCB AS&T)
June 30, 2010





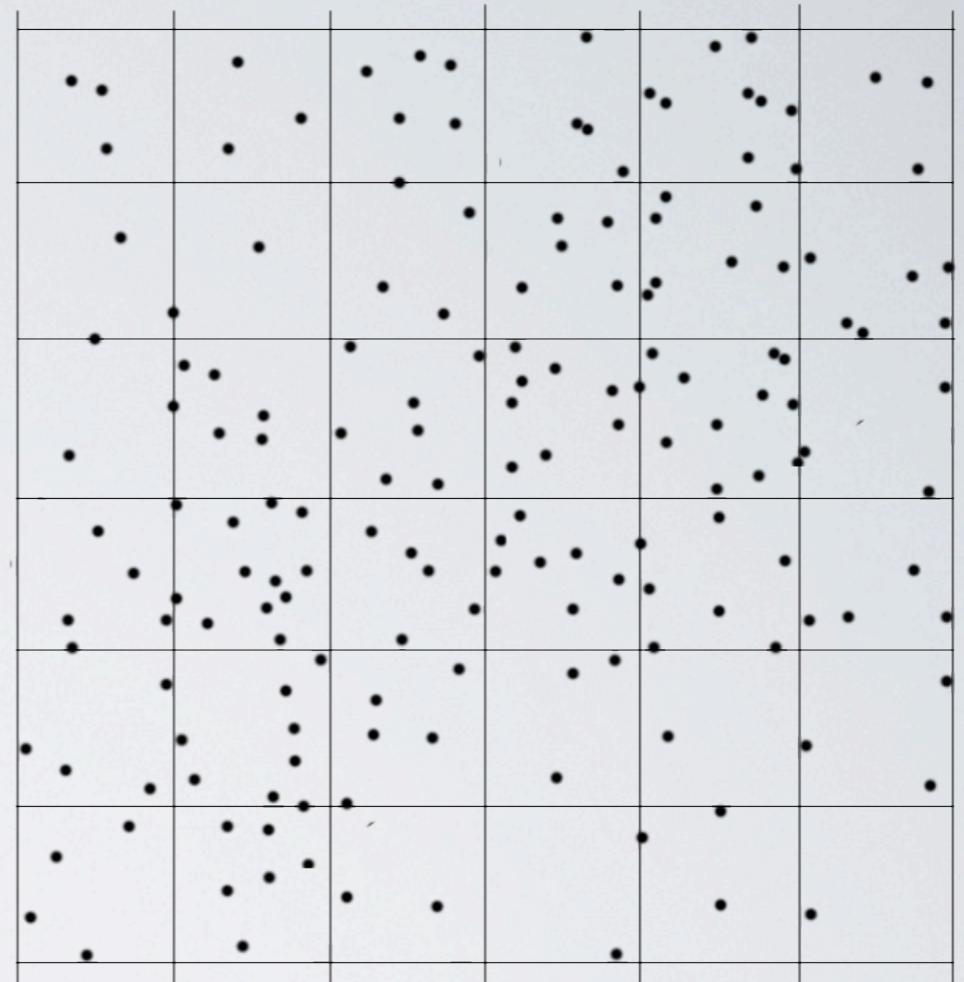
OUTLINE

- What is PIC
- Motivations for the Project
- Design Goals of Python PIC
- Interface Design
- Use Cases
- PyCUDA Test Problem and Results

WHAT IS PIC

$$\nabla^2 \Phi_{[i,j]} = -\frac{\rho_{[i,j]}}{\epsilon} \quad \ddot{x}_{(\vec{r})} = -\frac{q}{m} \nabla \Phi_{(\vec{r})}$$

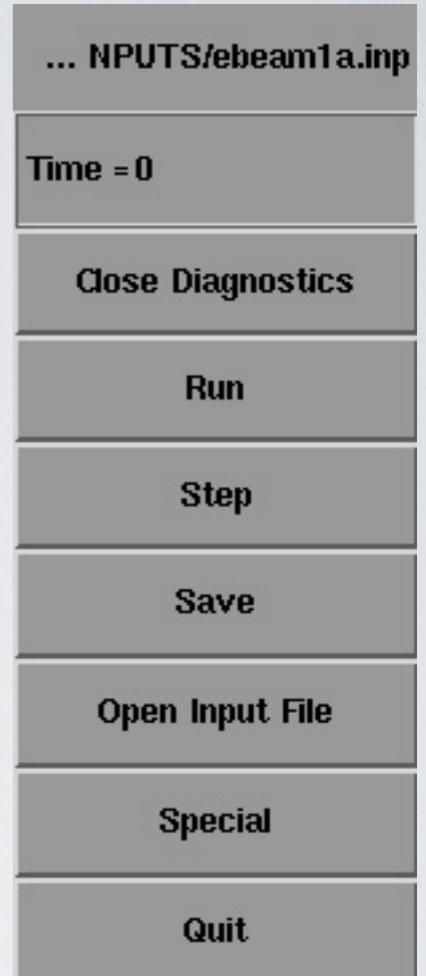
- Particle-In-Cell (PIC) Simulations are popular in plasma physics
- E/M Fields are defined on a discrete grid
- Particles live in continuous space
- Particles do not directly interact, but rather are mediated by the grid
- PIC:Weight (p->g), Solve(g), Interpolate (g->p), Push(p)





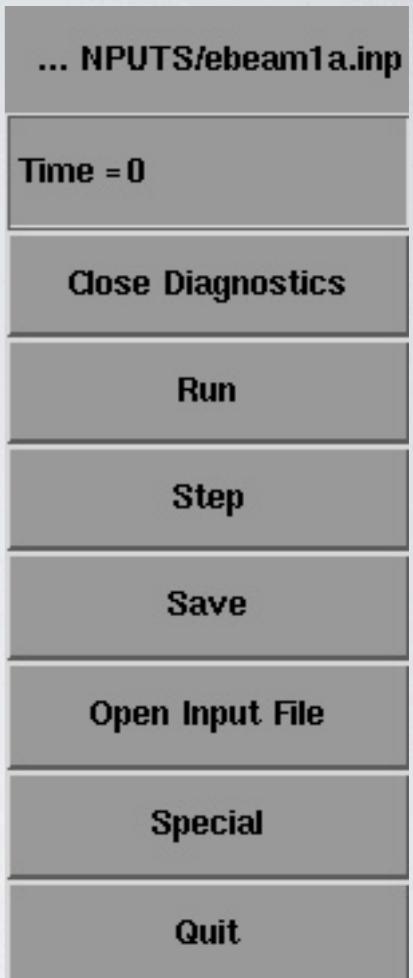
THE PAST: XOOPIIC

- Object Oriented Particle-In-Cell
- Rebuild the application if you add anything
- Good model, but not scalable
- Mouse-only interface to simulations
- Better tools exist now, that didn't in 1995



INTERFACE GOALS

- Represent Device as Python object
- Provide superset of XGrafix controls as methods of that object
- Present existing diagnostics as numpy arrays
- Allow arbitrary derivative diagnostics
- Pure Python for input files (no custom syntax)
- Multiple Backends (PyCUDA, 2D, 3D, etc.)
- Cluster management for parallel simulations (IPython)





PERFORMANCE GOALS

- Exploit ubiquitous parallel hardware to study larger problems
- GPU parallelism
- Multicore/Multi-chip parallelism
- Cluster parallelism
- Short Term: build a model that will use any of these
- Long Term: use all of them

INTERFACE DESIGN

INTERFACES

- Each type of object provides an Interface
- For extended functionality, subclass an Interface
- All Diagnostics provide `IDiagnostic`
- All Boundary Diagnostics provide `IDiagnostic` AND `IBoundaryDiagnostic`

```
    IDevice
    ↴
    IBackend
        ↴
        IDecomposedBackend
    ↴
    IDiagnostic
        ↴
        IBoundaryDiagnostic
        ↴
        IParticleDiagnostic
        ↴
        IFIELDDiagnostic
        ↴
        ITIMEDiagnostic
    ↴
    ISpatialRegion
    ↴
    ↴
    ↴
    IFieldSolve
        ↴
        ISolvePoisson
        ↴
        ISolveMaxwell
    ↴
    IParticleBin
    ↴
    ISpatialRegion
```

INTERFACES

```
import zope.interface as zi

class IDiagnostic(zi.Interface):
    data = Attribute("The Data for this diagnostic")
    interval = Attribute("The update interval")

    def save(fname=None):
        """save me to a file"""
    ...

class BaseDiagnostic(object):
    zi.implements(IDiagnostic)

    @property
    def interval(self):
        return self._interval

    @interval.setter
    def interval(self, newinterval):
        self._interval=newinterval
        self.backend.change_diag_interval(self, newinterval)

    def save(self, fname):
        self.data.tofile(fname)
```



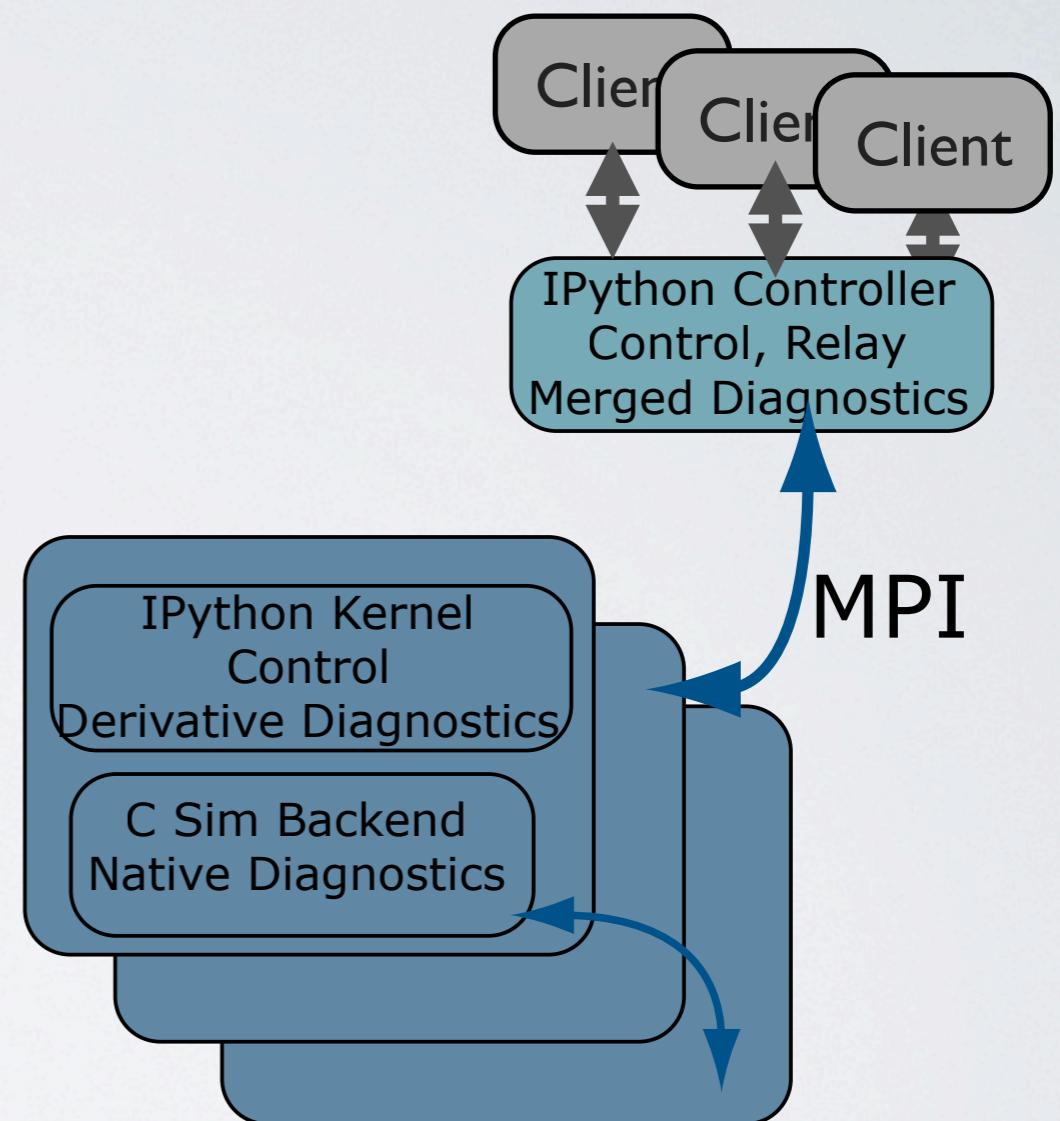
INTERFACES

- Rather than writing a parser, input files are Python scripts, executed in a namespace.
- That namespace is scanned for **Objects** providing **Interfaces**, and a Simulation is built from them.

THE DESIGN

USE MODEL

- User builds Logical Model in Python
- Python Backend builds Native Simulation Objects (C/CUDA)
- Control and Derivative Diagnostics in Python
- Inner Loops and Native Diagnostics in C/CUDA





WORKFLOW

- Solve Field
 - Compute Field Diagnostics
 - Push Particles
 - Particle Diagnostics
 - Weight Particles
 - defer to Python for Control
- Backend**
- Collect Derivative Diagnostic dependencies
 - Make adjustments
 - defer to parent for Control
 - <resume worker>
 - Compute Derivative Diagnostics
- Python**
- start next step(s)



PROXY OBJECTS

- In parallel Sims: diagnostics, etc. will actually be Proxy Objects, relaying instructions to Spatial Regions
- Data will not be retrieved unless explicitly requested
- Due to data moving, diagnostics will have update intervals at each level



NUMPY DIAGNOSTICS

- The basic data structure of all Diagnostics will be NumPy Arrays.
- This gives us (and users) free access to:
 - slicing
 - native BLAS,FFT,etc.
 - rich analysis tools with SciPy etc.
 - all major Python plotting tools - Matplotlib, Gnuplot, Chaco, etc.

USE CASES



```
from OOPIC import *
sim = Simulation("/path/to/inputfile.inp")
defaultDiags = sim.diagnostics
sim.run() # run simulation until interrupt
sim.run(10) # run 10 steps
sim.step() # step forward once
J = sim.diagnostics["Current"]
B = sim.diagnostics["Magnetic Field"]
JcrossB = cross(J[:,6], B[:,6]) # adds JxB Diagnostic
JcrossB.interval = 10 # update every 10 timesteps
pylab.plot(sim.X, JcrossB)

def f(N):
    return N > 1e9
n = sim.diagnostics["Number"]
sim.runUntil(f, (n,)) # runs until there are 1e9 particles
```



MULTIPLE SIMULATIONS

- As long as you can define your analysis, you can do it.
- Second simulation depends on results of first (sim2 could be DSMC or anybody else's code)

```
from mystuff import *  
  
sim1 = Simulation("sim1.py")  
def stop_condition(sim):  
    metric = check_some_stats(sim.diagnostics)  
    if metric > threshold:  
        return True  
sim1.run_until(stop_condition, (sim1,))  
  
analysis = perform_analysis(sim1)  
del sim1  
sim2 = Simulation("sim2.py")  
update_with_analysis(sim2, analysis)  
sim2.run()
```

SEARCH PARAMETER SPACE

- here we have a set of parameters that we want to find
- run a simulation, and get new input parameters from analysis
- start again with new parameters
- repeat until new parameters are close enough to the most recent input.

```
sim = Simulation("sim0.py")  
prev = None  
next = init_parameters()  
while prev is not None and \  
    norm(next-prev) > tol:  
    sim.reset()  
    configure(sim,next)  
    sim.run_until(stop_condition)  
    prev = next  
    next = analyze(sim)
```



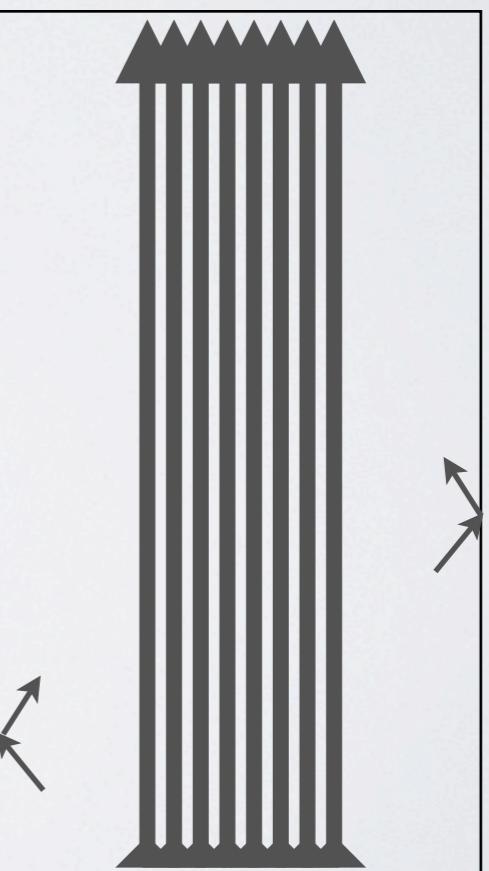
SUMMARY

- The UI and the Analysis systems are *the same*.
- Interfaced Based Backend model facilitates extensions of Physics
- Python provides intuitive, extensive Sim building tools
- Ground up Parallel design will scale better than XOOPIIC
- PyCUDA is the first Backend

THE TEST CASE

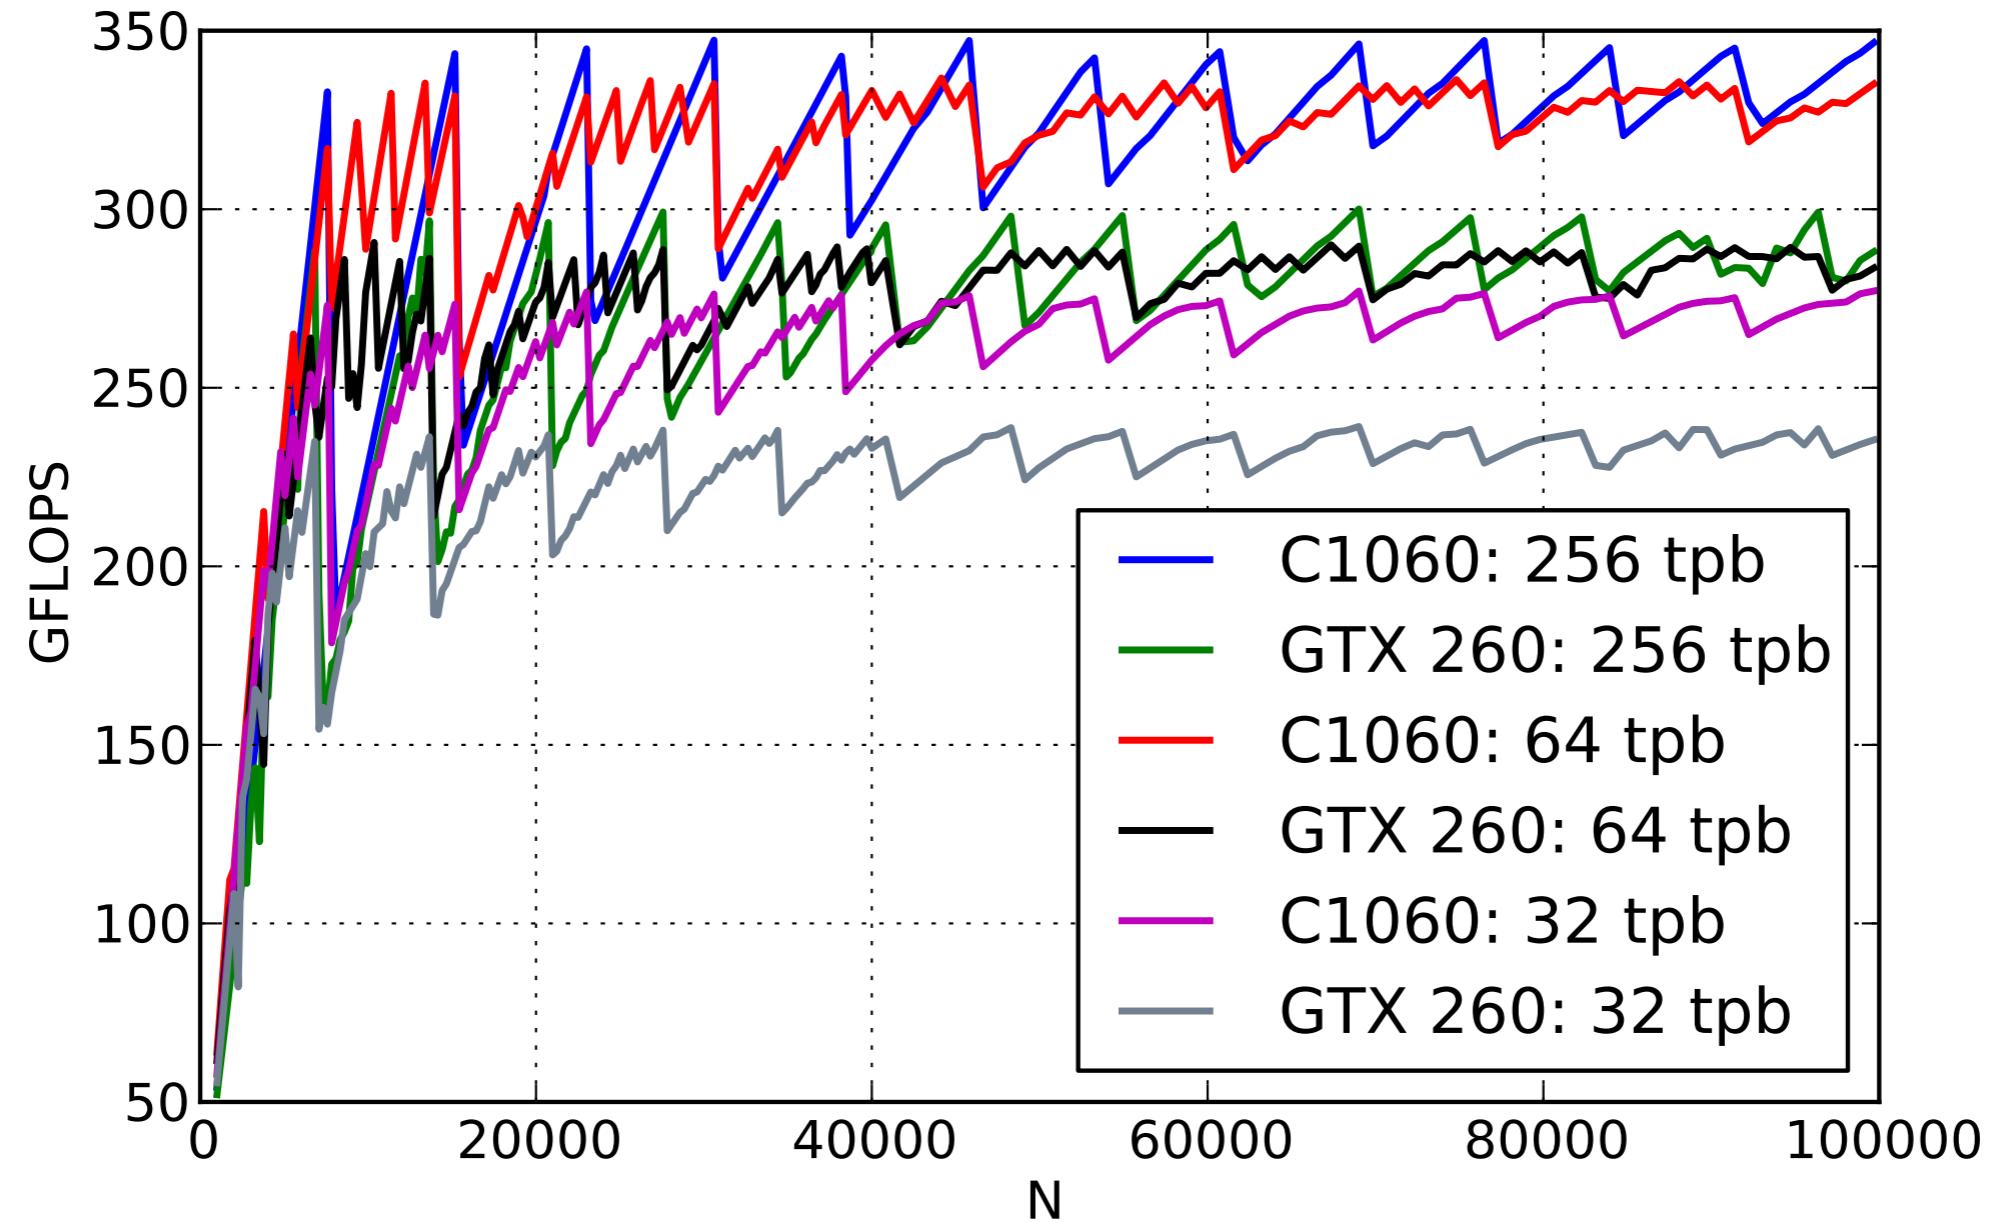
TEST CASE: SHEET BEAM

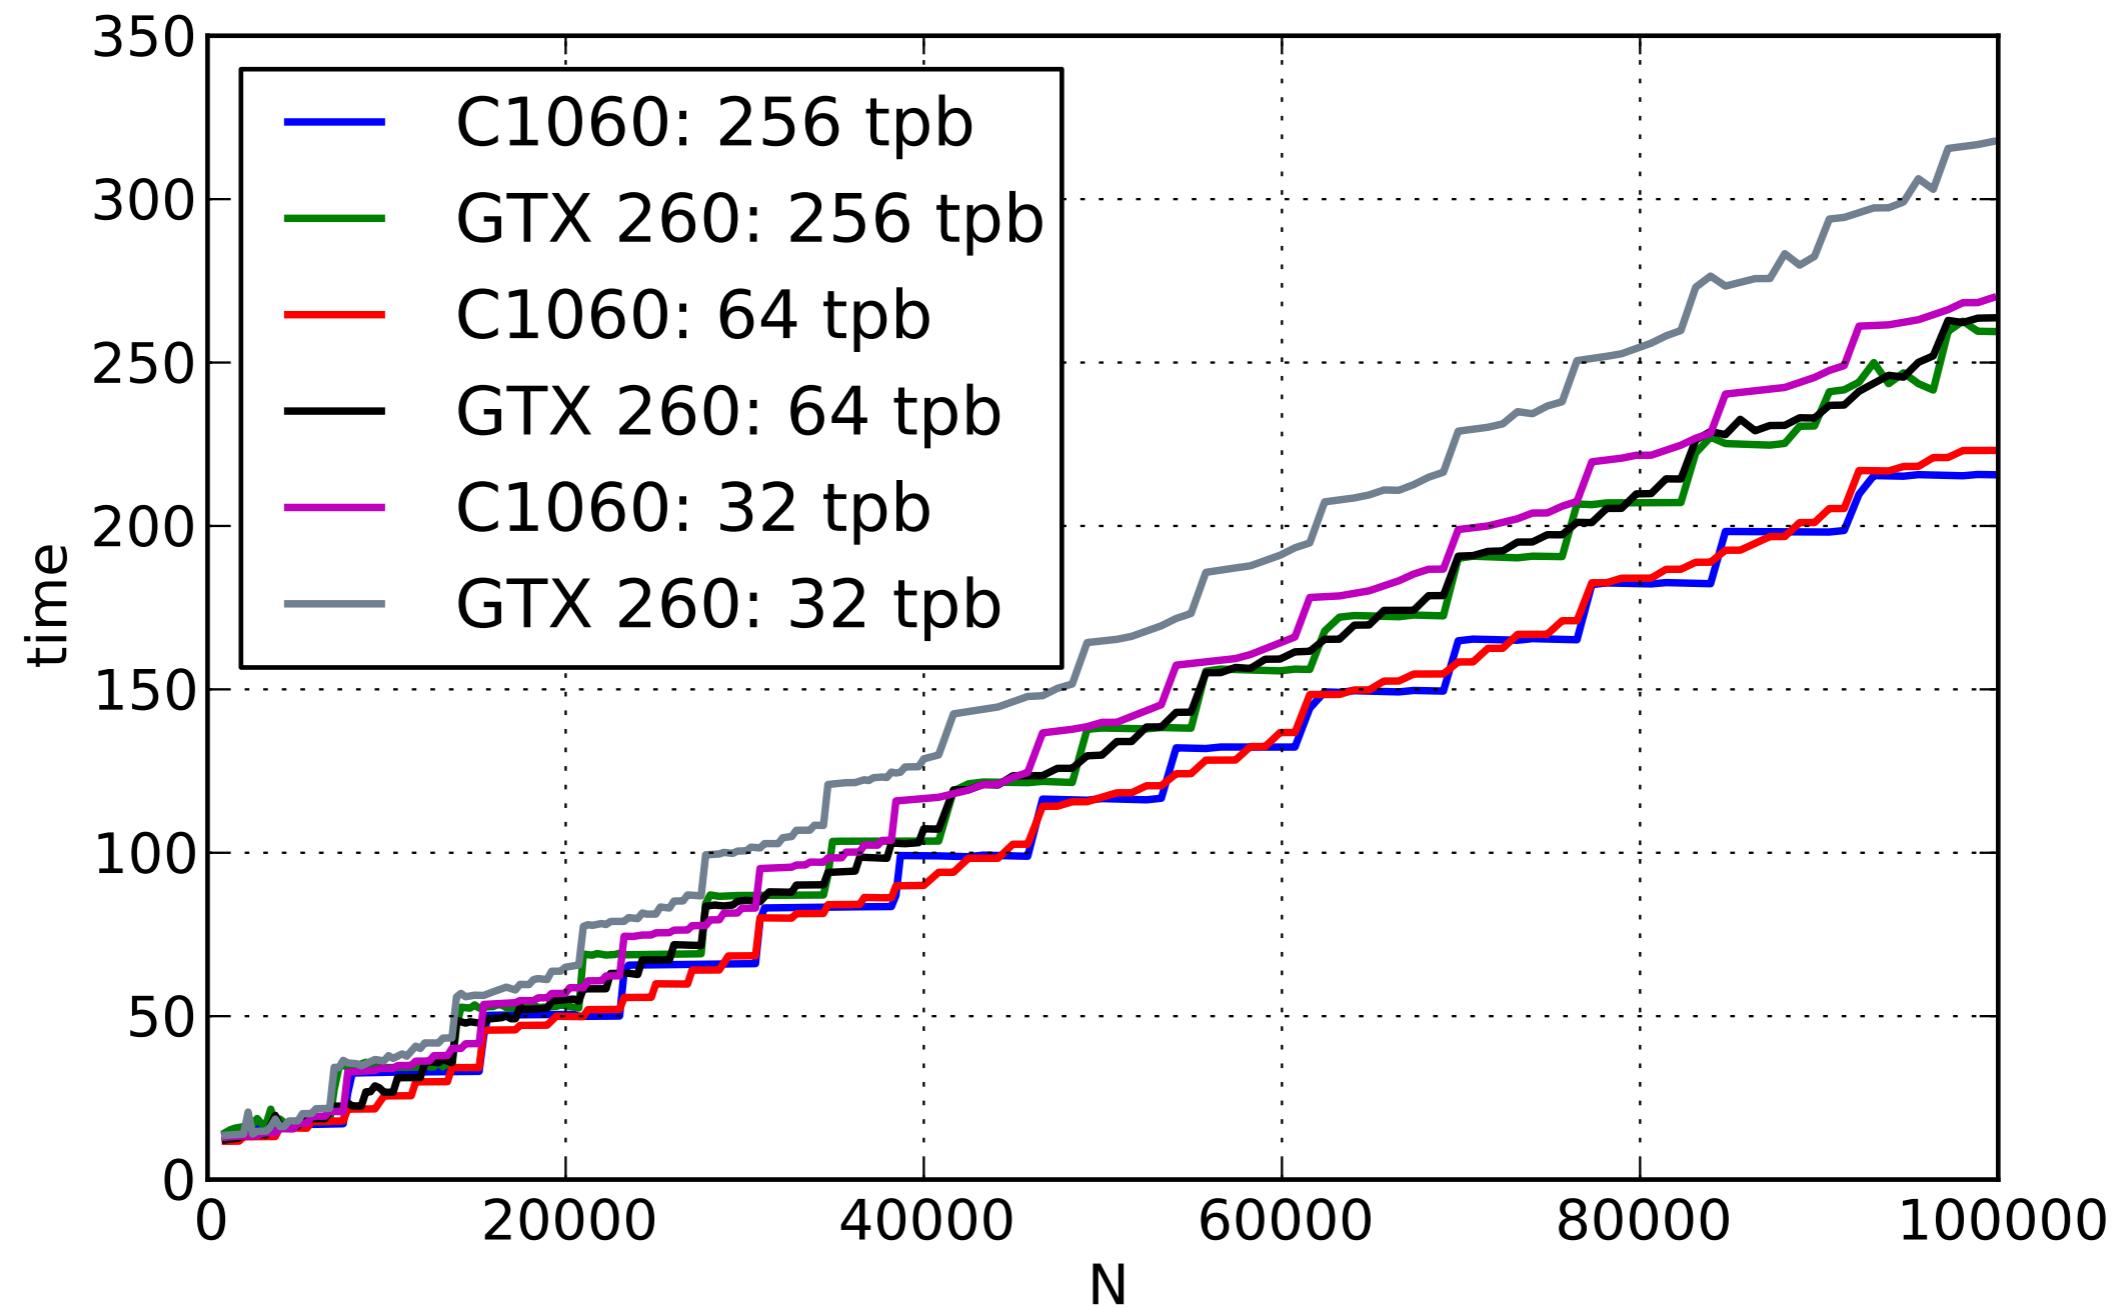
- Goal: Simulate a number of particles with PyCUDA (2D Cartesian)
- Direct Coulomb repulsion (short-range approximation)
- Initial Drift in +Y
- Specular bounce off X walls
- Periodic in Y (continuous infinite beam)

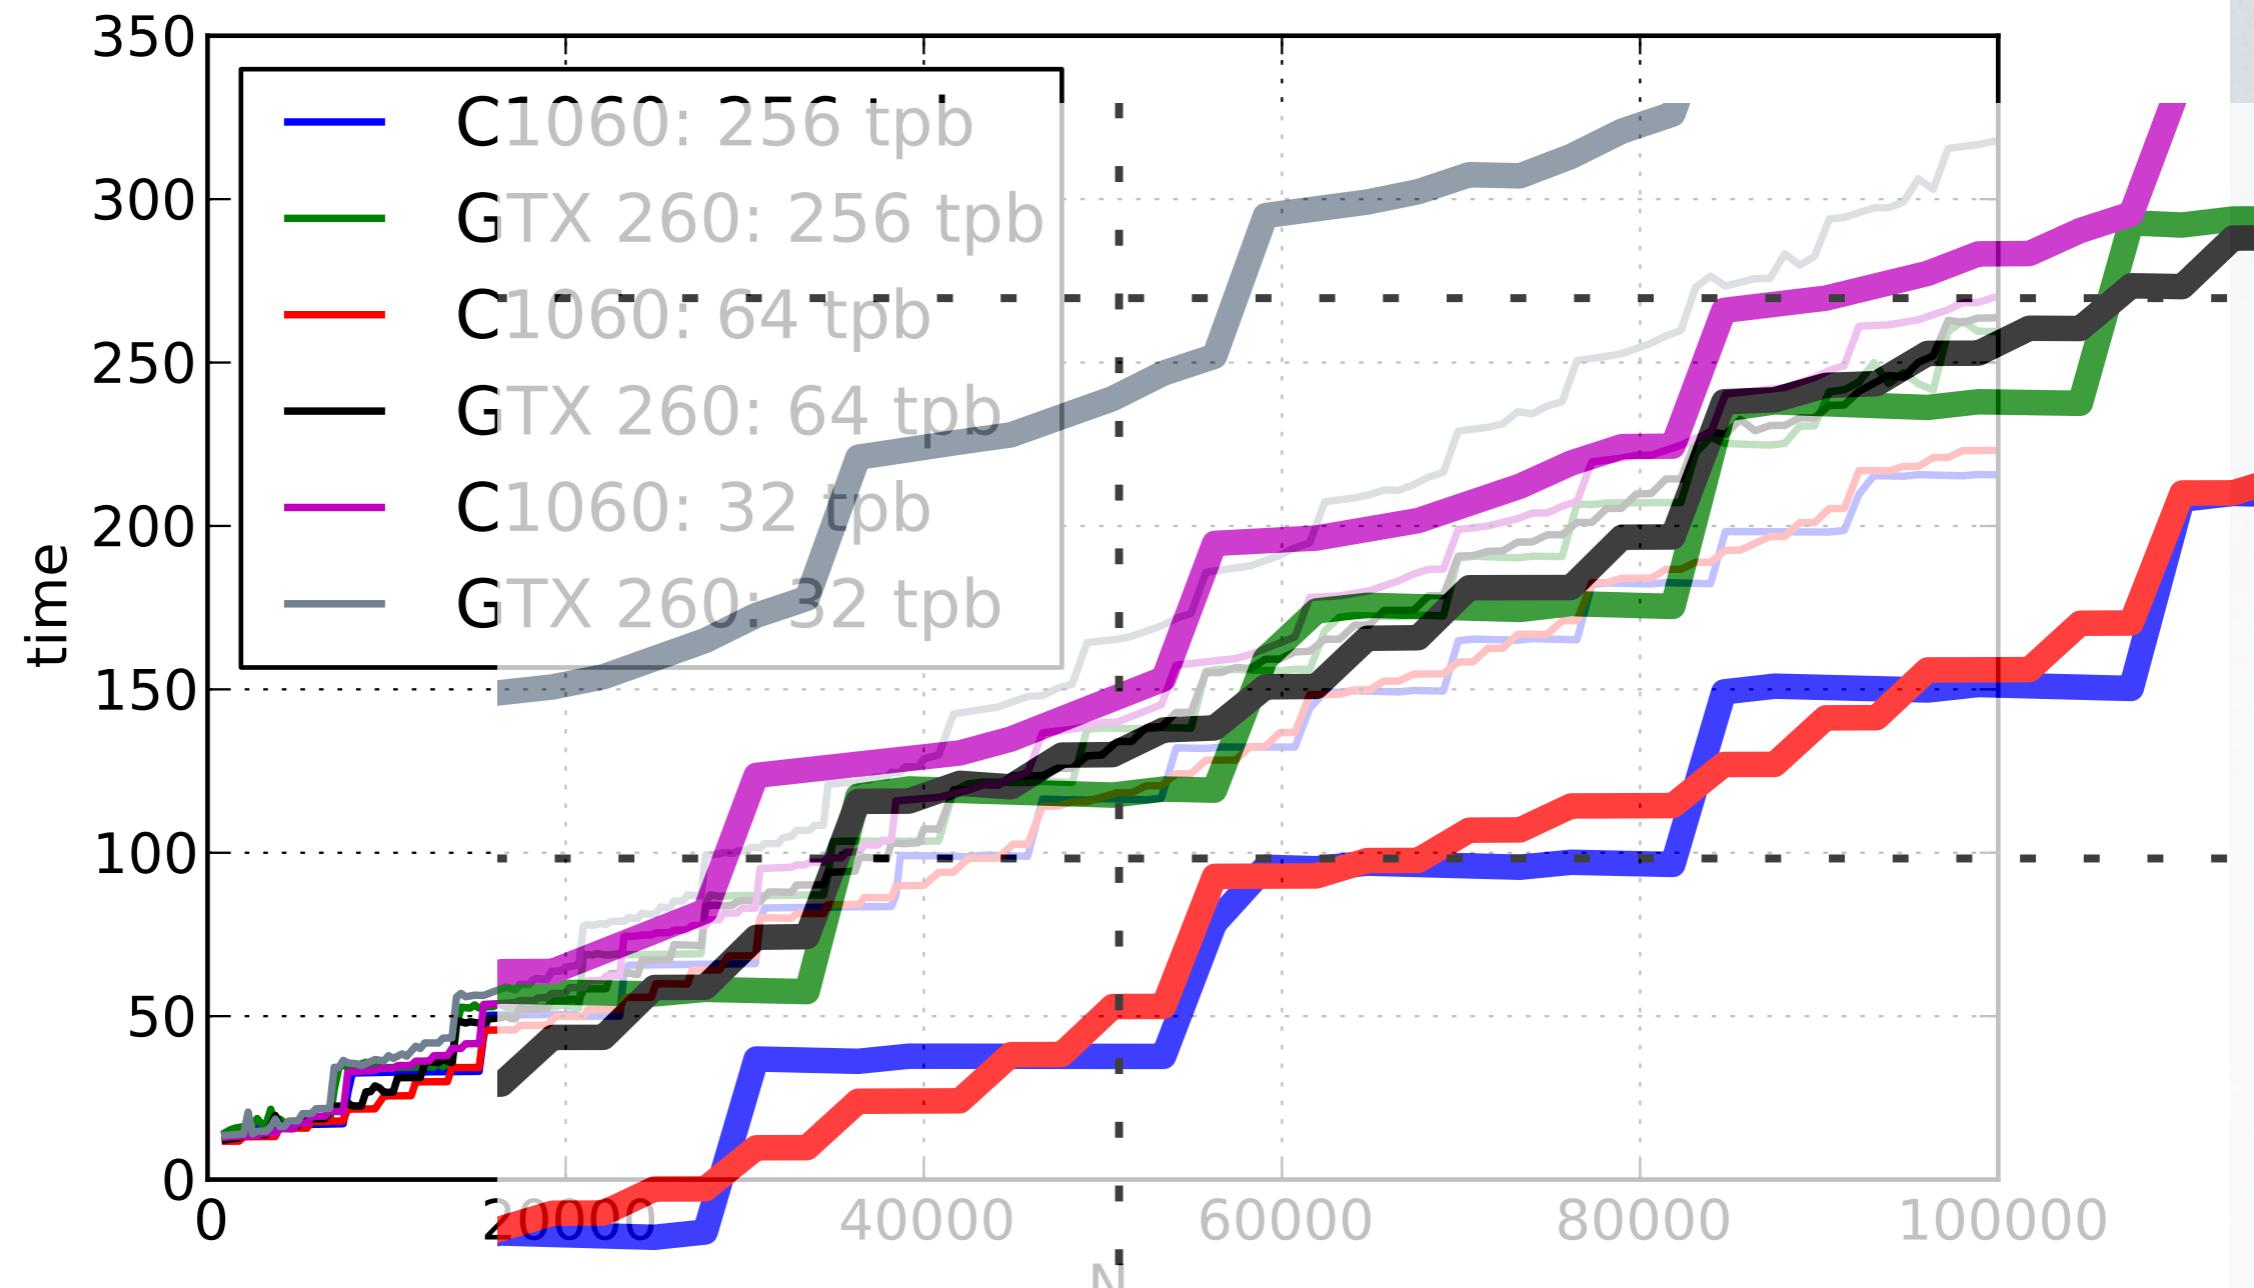


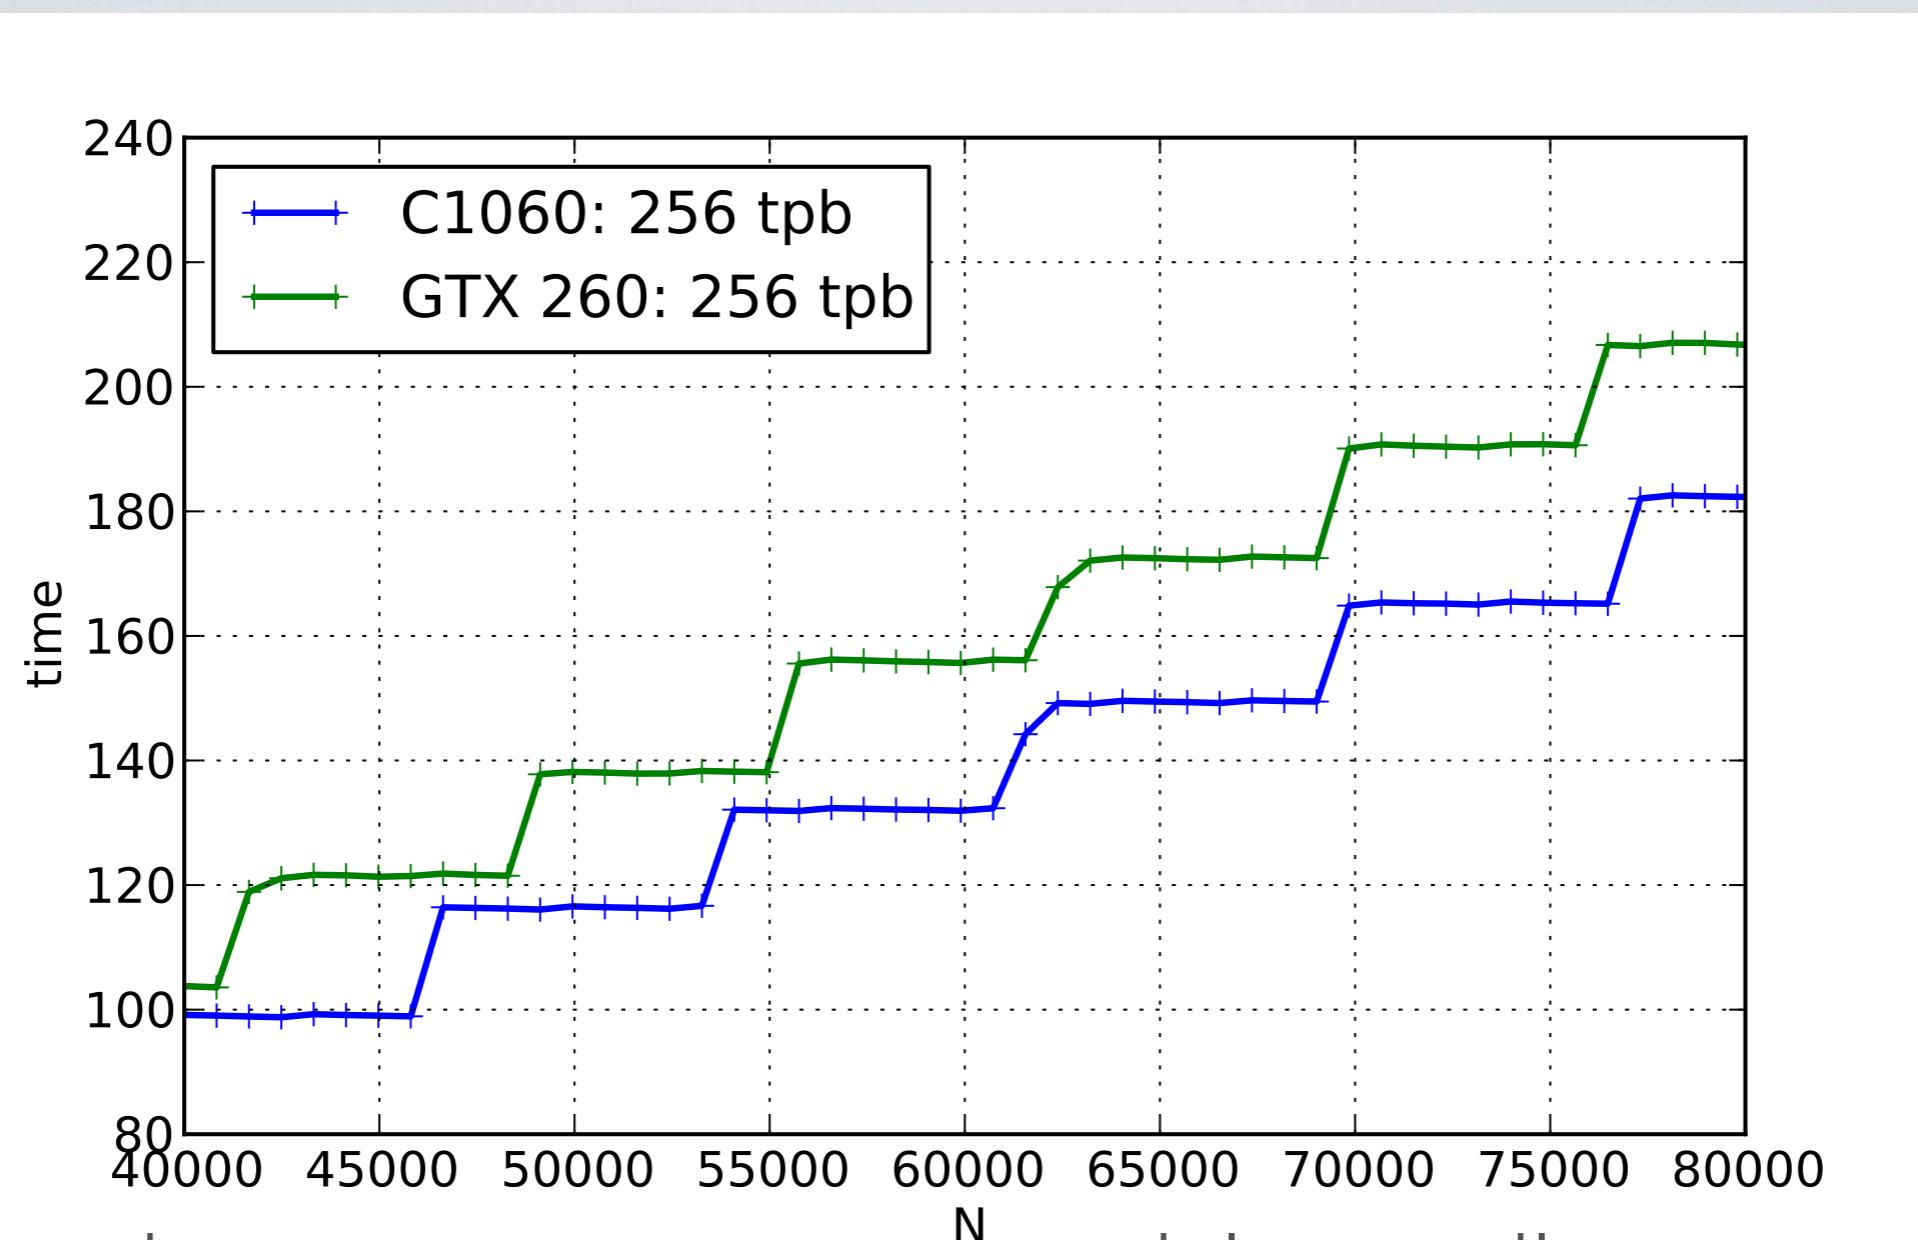
RESULTS

Performance

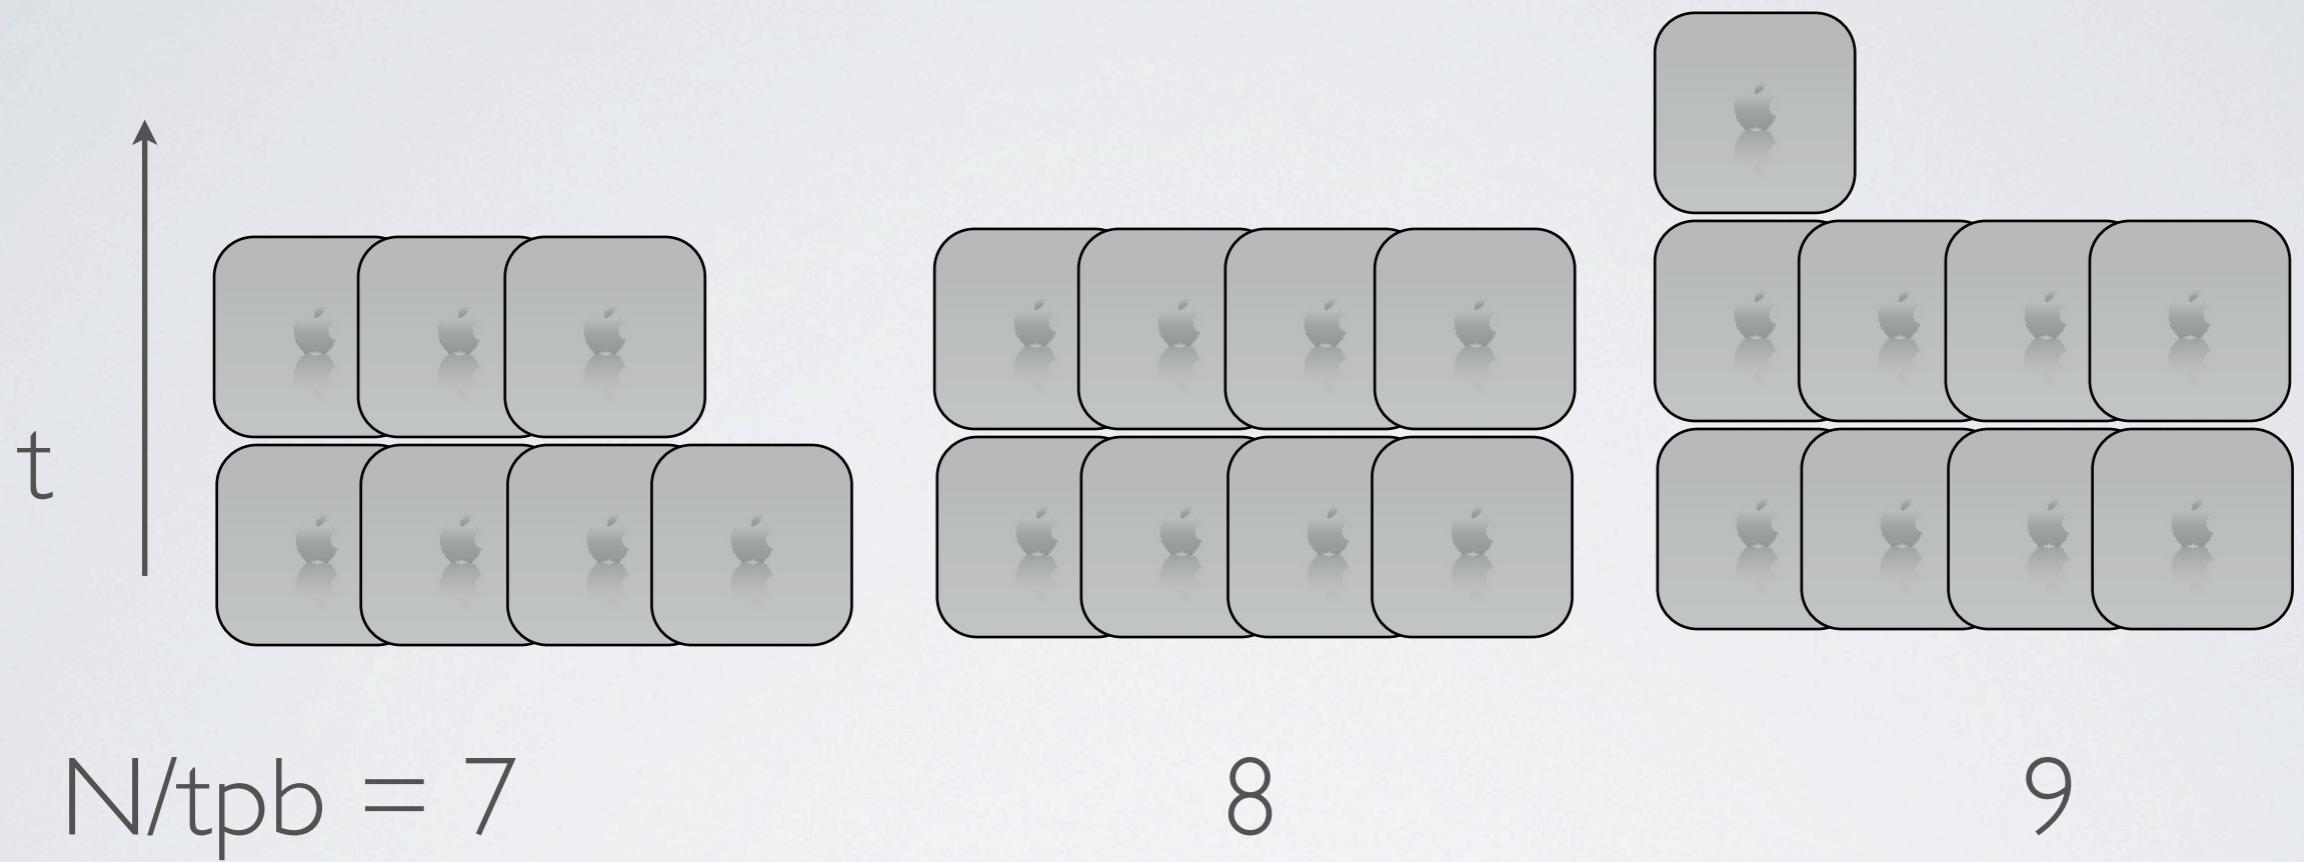


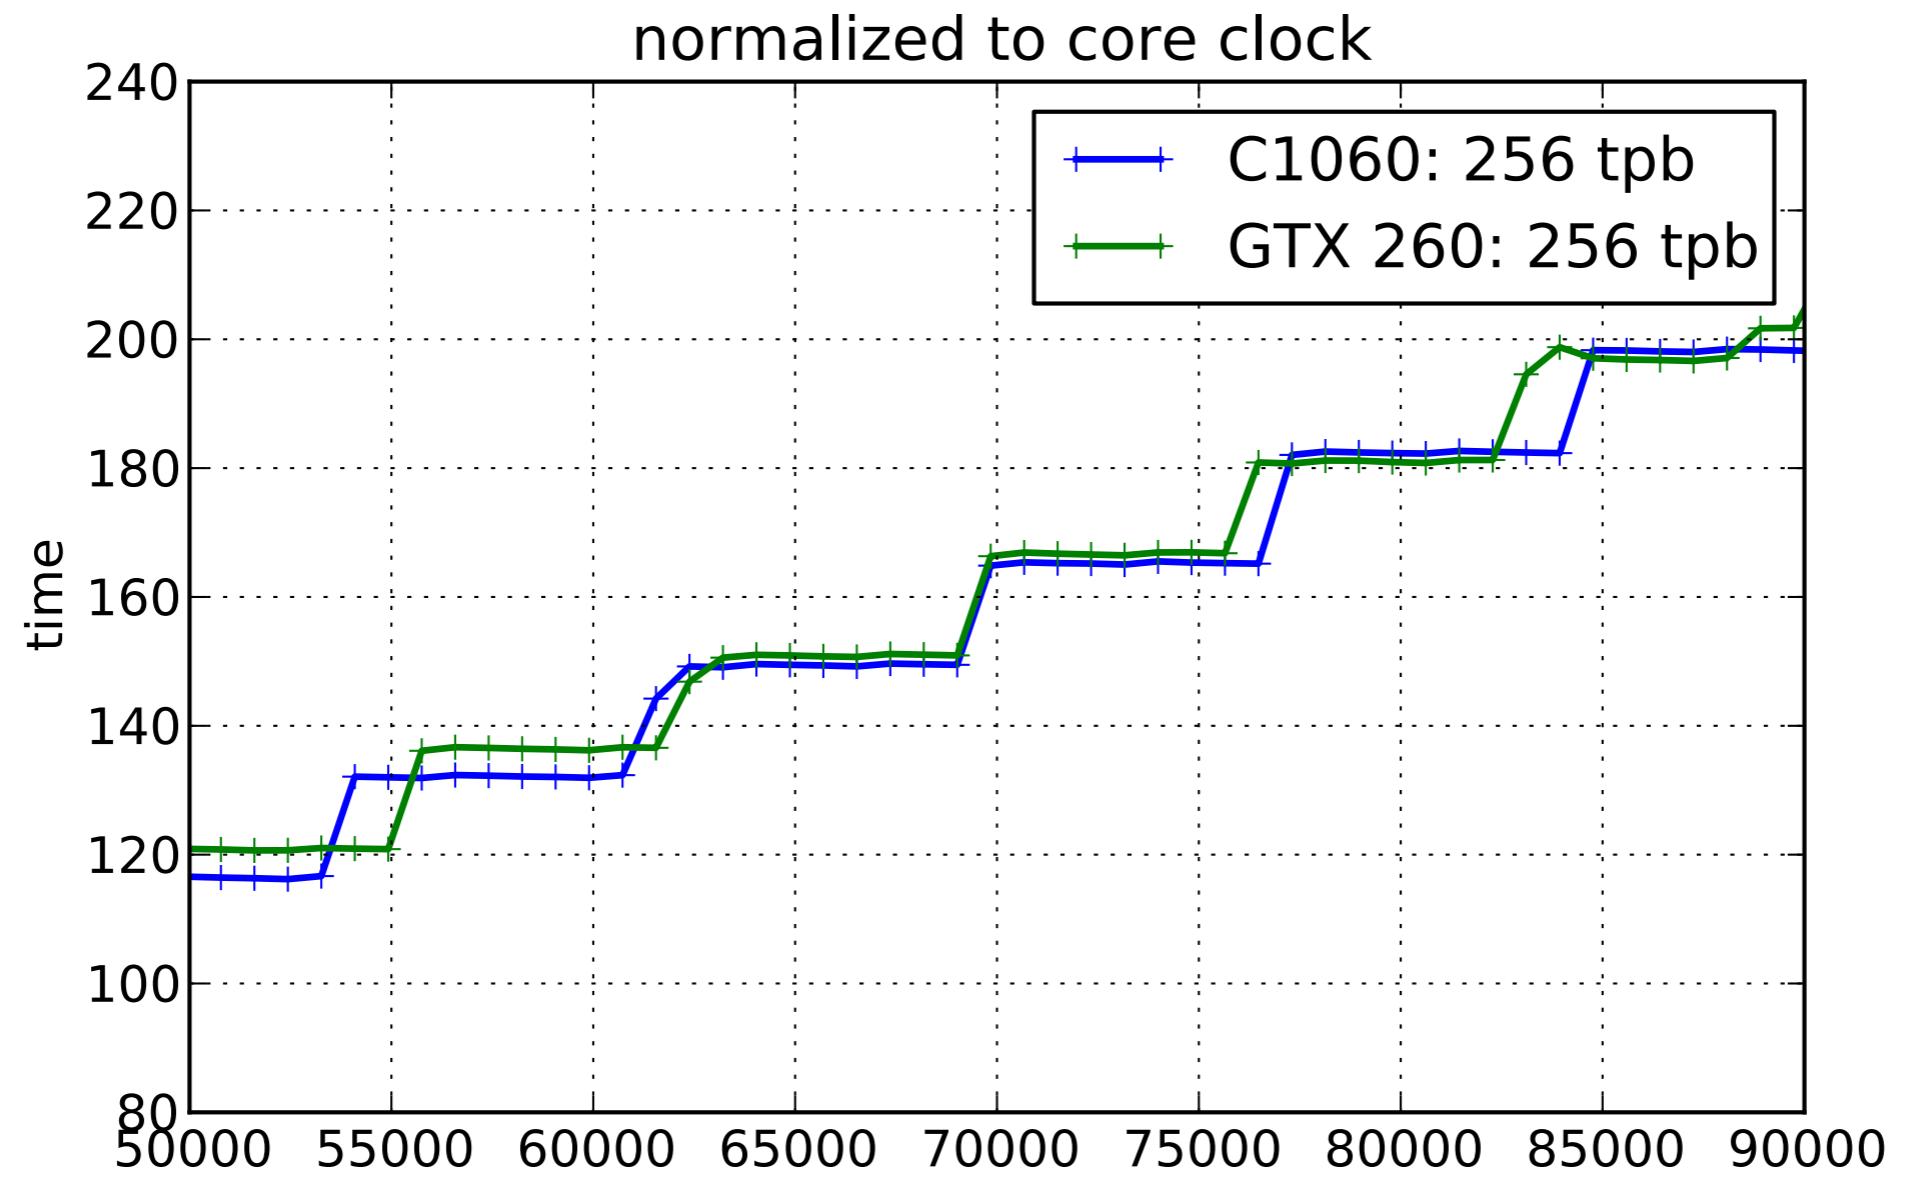






Zero time spent on new particles until a new set of blocks is needed





Adjusting for the difference in core clock,
 N performance of the two chips is very similar.



RESULTS

- Performance is periodic in (threads/block * GPU cores)
- The GPU is most fully loaded when the work can be broken into full-sized work groups
- when the work doesn't fit, some of the time is spent with the GPU not fully utilized

RESULTS

- Did achieve good (30% of ITFlop peak for C1060) performance on simple test case
- Did not get as far in exploring the problem as originally intended (comparing multiple schemes)



NEXT STEPS

- Develop test code into full PIC Backend kernel
- Plug the PyCUDA simulation into the Python Diagnostics and steering part of the code
- Eventually: Port PyCUDA backend to PyOpenCL in order to use MultiCore CPU as well



SOURCES

- CS267
- ParLab Bootcamp videos and pdfs
<http://parlab.eecs.berkeley.edu/bootcampagenda>
Specifically Demmel, Catanzaro, Mattson
- NVIDIA GPU Gems
http://developer.nvidia.com/object/gpu_gems_home.html
- NVIDIA CUDA/OpenCL Docs
 - http://www.nvidia.com/object/cuda_develop.html
- HPG/SIGGRAPH 2009



THANKS

- Fernando Perez & Brian Granger (IPython)
- Andreas Klöckner (PyCUDA/PyOpenCL)
- John Verboncoeur (Advisor)
- You (for coming)